Signal, Image and Video Processing

Automatic analysis of broadcast football videos using contextual priors

--Manuscript Draft--

Manuscript Number:	SIVP-D-16-00054R1
Full Title:	Automatic analysis of broadcast football videos using contextual priors
Article Type:	Original Article
Funding Information:	
Abstract:	The presence of standard video editing practices in broadcast sports videos, like football, effectively mean that such videos have stronger contextual priors than most generic videos. In this paper, we show that such information can be harnessed for automatic analysis of sports videos. Specifically, given an input video, we output per frame information about camera angles and the events (goal, foul, etc.). Our main insight is that in the presence of temporal context (camera angles) for a video, the problem of event tagging (fouls, corners, goals etc.) can be cast as per frame multi-class classification problem. We show that even with simple classifiers like linear SVM, we get significant improvement in the event tagging task when contextual information is included. We present extensive results for 10 matches from the recently concluded football world cup, to demonstrate the effectiveness of our approach.
Corresponding Author:	Rahul Anand Sharma, MS International Institute of Information Technology Hyderabad INDIA
Corresponding Author Secondary Information:	
Corresponding Author's Institution:	International Institute of Information Technology Hyderabad
Corresponding Author's Secondary Institution:	
First Author:	Rahul Anand Sharma, MS
First Author Secondary Information:	
Order of Authors:	Rahul Anand Sharma, MS
	Vineet Gandhi, Ph.D
	Visesh Chari
	C. V. Jawahar
Order of Authors Secondary Information:	
Author Comments:	
Response to Reviewers:	We thank reviewers for their comments and insights. Below are the answers to their queries: Answers R1-The manuscript is difficult to read and understand. References are not orderly cited.... Ans. References are ordered automatically by the provided latex template for the journal. We have incorporated other specific comments and have tried to improve the readability. R1- Details of the methods are not clear. They should be explained in detail, especially for frame representation and cost function. Ans.Done. For representation, we have now cited paper which proposed the original bag of words approach, that the readers can refer for a detailed explanation. Cost function explanation has also been elaborated.

R1- What does N refer in equation 1? What is the value of Lambda in equation 2?
Ans. N is total number of the frames (explained in the updated manuscript).

R1- A few sentences about dynamic programming are necessary in terms of optimization problem.
Ans. Incorporated

R1- For classifications, why 40 frames are used? why 30 not? why 60 not?
Ans. We empirically tested with different numbers. Forty gave a good tradeoff between accuracy and computational load. We did not put this analysis due to lack of space (8 page limit).

R1- Please, define MRF as "Markov Random Field" if I am right.
Ans. Yes it is Markov Random Fields. We have updated this in the paper.

R1- Number of quantitative results are sufficient. But which are ground-truth data in Tables (from 1 to 6)? Columns or Rows?...
Ans. Ideally it should be diagonal (with all diagonal values equal to 100 and rest zero). Yes each row should sum to 100, there was typo during rounding, we have corrected that. We thank you for carefully reading and pointing this out.

R1- According to the manuscript, Two 45 minutes videos mean 324000 frames. Should it be 162000?
Ans. Videos are sampled at 60 fps. So two 45 minutes is equal to 2*45*60*60=324000

R1- Number of qualitative results are sufficient. But the results are not convincing. For example, if we look at the gameplay.mp4 video, we see the likelihood of the occurrence of an event at the current frame. Let's look at the second highest peak which is the goal event. The goal event must have very low likelihood according to the quantitative results. Similar negative conditions can be seen from the other videos.
Ans. The values in plots denote just the relative ordering (ranking) of events based on their predicted probability (the peaks are not showing the exact confidence score). So even though goal event is second highest peak in that graph, its confidence score predicted by the algorithm may be quite low (which is usually the case).

R2- In page 3, line 48, column 2; authors say that camera viewpoints can be broadly categorized into five different categories. I wonder if this analysis was empirically made or in a supervised manner...
Ans. This labeling is to cover all possible camera labels. We performed empirical analysis on several matches and found out that the first four labels (ground zoom-in, ground zoom-out, top zoom-in, top zoom-out) cover nearly 98% of the match. We classify the rest into miscellaneous. Because of the low occurrence of the Misc class, further subdivision is unnecessary.

R2: In Table 1-2-3, authors demonstrate that their method performs better in camera label estimation compared to dominant color ratio and optical flow. For camera label estimation, their method utilizes temporal classification data using a window of 40 frames. Is this framework also applied to dominant color and optical-flow results? If not, this may not be a fair comparison.

Ans. Yes, the temporal window of 40 frames was applied in other cases as well.

R2: Also in Table 1-2-3, is it possible to provide execution times, including feature extraction process. It would be nice to see the trade off between time and accuracy.
Ans. Our algorithm for camera label estimation runs nearly in real time (20fps in current implementation). The dominant color based approach is way faster (nearly 60-100fps). The motion cue based approach

R2: Again in Table 4-5-6, it would be nice to include the execution time or complexity of different steps. For instance, adding motion features in the last step increases gameplay classification results by 3.2%. What is the cost of this?
Ans. We have compared and updated the paper accordingly.

R2: Finally, and most importantly, authors only verbally compare their approach with

	the related studies. I thing quantitative experiments are lacking in highlighting their contributions compared to previous work. I would like to see a comparison with the direct results of state-of-the art approaches for the results they provided in event tagging. It could be possible to find and access results of related studies for events such as goals. Ans. Most of the previous approaches are based on ad-hoc rules which makes it difficult to replicate their algorithms. Their codes and the dataset are also not available, so making a fair comparison is not trivial. Even on a quick glance, it is easy to observe many of the adhoc rules used in previous methods won't apply in the challenging data-set proposed in our work. For instance, many previous methods are based on hough line detection, which completely fails in presence of shadows or motion blur etc. Many others use the optical character recognition on scorecard (whose location is hard coded and known). Some use logo detection, or replay detection or audio cues. All such cues vary across broadcasters, tournaments etc. and are extremely difficult to replicate (lot of manual rules or labeling is required to obtain such cues). But since you have asked we have compared our method with a generic deep learning based state of the art action recognition approach and a mid level semantics based sequence classification technique. The results have been updated in the manuscript. We also implemented and tested few other rule based goal detection methods and found that their results were much inferior and not appropriate to be included in the manuscript.

Noname manuscript No.
(will be inserted by the editor)

Automatic analysis of broadcast football videos using contextual priors

Rahul Anand Sharma · Vineet Gandhi · Visesh Chari · C V Jawahar

Received: date / Accepted: date

Abstract The presence of standard video editing practices in broadcast sports videos, like football, effectively mean that such videos have stronger contextual priors than most generic videos. In this paper, we show that such information can be harnessed for automatic analysis of sports videos. Specifically, given an input video, we output per frame information about camera angles and the events (goal, foul, etc.). Our main insight is that in the presence of *temporal* context (camera angles) for a video, the problem of event tagging (fouls, corners, goals etc.) can be cast as *per frame* multi-class classification problem. We show that even with simple classifiers like linear SVM, we get significant improvement in the event tagging task when contextual information is included. We present extensive results for 10 matches from the recently concluded football world cup, to demonstrate the effectiveness of our approach.

Keywords sports video analysis · broadcast video · event classification · content based retrieval

1 Introduction

Modern developments in multimedia creation, storage, and compression technologies have paved the way for extensive archiving of video content. Building applications for search, summarization or editing on such large databases requires extensive information about the content of these videos. Current sources of such descriptions are only limited to textual content. Since textual descriptions are both inefficient (descriptions are subjective and vary from person to person) and incomplete (it is difficult to describe all content in a video to facilitate search, summarization or editing), it is important to build tools to automatically analyze video content and identify salient parts, to generate textual and other kinds of descriptions over the timeline. Given the diverse nature of video content on the web, this task is easier said than done. One approach to this problem is to isolate and process videos by genre. Such an approach has a two-fold advantage: 1) each genre could be associated with a set of rules for video creation that might make it easier to design video understanding algorithms 2) it is easier to distinguish between relevant and irrelevant semantic content when information about genre is given (for example, information about crowds in a football match is rarely searched for, and hence can be ignored). Recently, problems related to sports video analysis have particularly received much attention in this direction with many direct applications like automated highlights generation [1] or analysis of team activities and strategies[13], being built upon semantic analysis of video content.

Following these lines, we have looked at the problem of automatic semantic analysis of football broadcast videos. Our work is built on the fact that broadcast videos of football matches are constructed in a very structured manner, thus imposing some useful restrictions on the content. Firstly, there are fixed vantage points in a football field where PTZ cameras are placed for recording such events, thus limiting the number of views. Secondly, the edited video switches between these cameras in strong correlation with events occurring on the field.

This inherent structure of editing in broadcast sports videos motivates us to ask the question, how to define

R. Anand
Center for Visual Information Technology,
International Institute of Information Technology,
Gachibowli, Hyderabad - 500 032, Telangana, India
E-mail: rahul.anand@research.iiit.ac.in

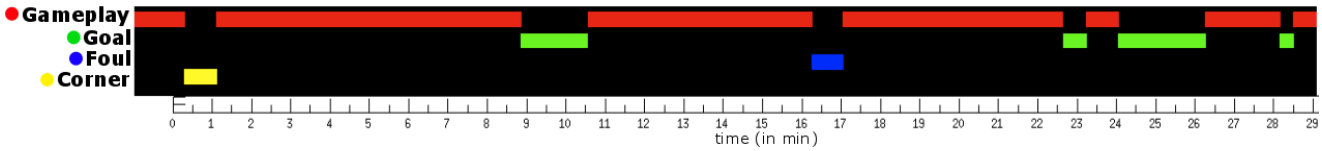

Fig. 1 An example of event tagging using the proposed approach on the first thirty minutes of the semifinal match between Brazil and Germany in World Cup 2014. The plot shows the occurrence of four different events (Goal, Foul, Corner and Gameplay) over the timeline. The proposed method successfully detects all the goal events.

context in such structured settings? We answer this problem based on two key realizations. Firstly, most camera angles associated with the events like goal or corners are pre-determined. For example, in the event of a goal, a broadcast video automatically switches to focus in on player huddles/celebrations, which is unlikely to happen during normal game play. Secondly, this strong association also extends temporally since editing rarely switches between unrelated camera angles, which is to say that camera angles vary *smoothly* across time, except when shot changes occur. This strong temporal and event correlation with camera angles prompts us to argue that *contextual* information in sports video analysis can largely be based upon knowledge of camera angles.

Based on the above discussion, we argue that for the problem of automated sports video analysis two main tasks gain prominence. Firstly, classification of images into different camera views forms the initial basis of *contextual* understanding of a sports video. Secondly, the analysis of *events* in such videos gives an almost complete summary of the match (a motivating example is illustrated in Figure 1). Accordingly, we focus on these two aspects in the current paper. Furthermore, as evidence of the usefulness of such information, we present an application of context reliant targeted spatial segmentation. Formally, we make the following contributions

- We present an automatic approach to first identify the camera view type for each frame in a video.
- Using predicted view types we then propose an algorithm to accurately predict salient events like goal, foul, corner, substitution etc.
- We also present an application of spatial segmentation that benefits from such contextual information.

2 Related Work

Previous work on semantic segmentation of videos has looked into two different directions, first is segmentation based on shots and second is based on events. Shot based temporal segmentation divides the video into smaller segments by identifying transitions from a camera to another. Most existing methods [11][18][14] performs shot segmentation by detecting transitions (cut, fade, dissolve etc) based on difference of features in consecutive frames. The segmented shots are then used as minimal units for further tasks.

In [25], Xu et al. classified football video shots into the views of global, zoom-in and close-up using the grass area ratio. In [24], Gu et al. used motion vectors in addition to dominant color to classify different view types in football broadcasts. Duan et al. [4] used a more elaborate set of features fusing motion vectors, color, texture, shape and shot length information for a similar task. These approaches suffer from the drawback of strong reliance on detecting transitions which may not be robust due to strong correlation between different shots (for example, a camera change may not result into difference in frame color histograms which are commonly used features in these approaches). To suppress the negative effects of inaccurate shot boundaries, it is also common to manually label the transition frames [25]. Our work addresses this issue in a more elaborate manner, by first training detectors for different camera types (illustrated in Figure 2) and then assigning each frame an unique camera label. By merging consistent camera labels over time our method in turn can be used to obtain shot based segmentation.

On the other hand, event based segmentation divides the video into shorter clips where each clip contains only one event. The granularity of event is lower than that of shot (an event can compass several shots) and there is no standard relation between an event and the number of shots in it. In [17], Qian et al. define an event clip as a set of sequential video frames which begins with a global view (far zoom out view) and ends with non-global views. They further use hidden conditional random fields to classify event clips into five different categories like goal, shoot, normal etc. But such a hard coded definition of an event may not always hold. Sigari et al. [18] segment the video based on camera motion (estimated using block matching between consecutive frames) and uses this information to detect event like counter-attack. Xie et al. [22] used a sliding window approach to classifying the football video into play vs break segments.

Fig. 2 We classify camera viewpoints into five different categories namely (from left to right) ground zoom-in, ground zoom-out, top zoom-in, top zoom-out and miscellaneous (covering mainly the crowd view).

Heuristic rule based approaches [22][1] have also been proposed for the detection of event boundaries. These approaches perform low level analysis to detect marks (field, lines, arcs, and goalmouth), player positions, ball position etc. A set of hard coded rules based on these low level features are then used to detect corresponding events. Assfalg et al. [1] used such a rule based approach for identification of salient events like goal, corner, kick off etc.

Recently purely data driven approaches have gained importance in sports broadcast, for example [2] learnt directorial styles by training classifiers on a training set of previous broadcasts. They suggest that such an approach could also be useful to determine the boundary of the salient events. Following these line, we employ a simple bag-of-features [3] representation combined with Support Vector Machines (SVM) for both view-type and event classification. The advantage of our approach is that it is independent of any ad-hoc rules or hard coded definitions and thus it is more generalized and can be easily extended to different sports.

Multimodel approaches are also common. Previous work [8] has looked at the problem of event detection using fusion of audio visual features. The main idea in these approaches is that the increased audio activity is a cue for important moments in the game. Textual information has also been used with visual features for event detection [23]. These methods typically use the match report and game log obtained from the web as the text source and require the alignment of the web-casting text with the broadcast sports video. On the contrary, the proposed method in this paper is purely based on visual data.

Recent work [13] has looked at the problem of discovering team behaviours or detecting the locations where the play evolution will proceed [9] by analyzing plan view tracks of the players. The plan view tracks are usually obtained using a set of static cameras manually installed around the corresponding sports field. Although these methods have demonstrated impressive results, they are not applicable for analysis in broadcast videos where only the feed from a single moving camera is available at a given time.

Far zoomed out shots have been used for the context of sports classification [7] and recognition of group activities in football videos [10]. The work in [7] exploits the fact that the playing surface is largely visible in far zoomed out viewpoints and different sports can be classified based on the type of playing surface. In [10] Kong et al. built a local motion descriptor by grouping SIFT keypoint matches into foreground point set and background point set and then used it to classify events like left side attacking or left side defending etc. Applicability to only far zoomed out shots limits these approaches in the case of broadcast videos which require ability to process input from different viewpoints (as illustrated in Figure 2). Our approach, on the other hand takes advantage of different view-points and presents a comprehensive solution.

3 Method

In this section we discuss the two key components of the proposed method for automatic analysis of football videos, namely, camera viewpoint estimation for each frame of the video and marking salient events on the timeline (event tagging).

3.1 Camera viewpoint estimation

Football broadcast videos are usually edited from a set of source videos recorded from different viewpoints. By analysing 52 matches of world cup 2014, we found out that these camera viewpoints can be broadly categorised into five different categories, which are illustrated in Figure 2). The inclusion of camera angles (ground or top) accounts for the main difference over the previous works, which segregate the camera view points only on the basis of their shot sizes (for e.g. medium shot or a close up). Although most of the game play is covered by top view cameras located in the crowd area, the ground view cameras (placed at sidelines at field level) plays an important role in viewing experience. They are quite handy during breaks and salient events and are used as standard by all most all broadcasters. An example of camera switches during a goal event is illustrated in Figure 3.

Our goal is to automatically predict the camera viewpoint for each frame of an edited football broad-

Fig. 3 Typical viewpoint transitions in a goal event (Top zoom out → Top Zoom in → Ground Zoom in → Ground Zoom out → Miscellaneous). The camera framing changes with respect to both size and angle.

cast video, as this provides a strong context for other higher level tasks (for example, the pattern of transitions between camera viewpoints is correlated with different salient events as illustrated in Figure 3). We approach this problem by learning a per frame multi-class classifier.

Frame representation: Each frame is represented using the classical Bag of Words (BoW) approach [3]. SIFT features are first computed for each frame independently, which are then clustered using k-means clustering algorithm to build a visual vocabulary (where each cluster corresponds to a visual word). Each frame is then represented by the normalized count of number of SIFT features assigned to each cluster (BoW histograms). The length of the feature vector is equal to the number of clusters.

Using the BoW feature representation, our method learns a multi class SVM to classify different camera labels. Training is performed by manually annotating a set of frames with corresponding camera labels. The classification is performed per frame basis but for classifying a frame we consider features from a temporal window of 40 frames centred around it.

We further assume that the camera transitions (cuts) are smooth and each camera viewpoint is maintained for a minimum amount of time before cutting to a different camera. This is a fair assumption, as fast and abrupt camera transitions may appear disturbing to the viewer and are avoided by the expert editors [19]. We benefit from this assumption to smooth out the noise in per frame camera label prediction, using a Markov Random Field (MRF) optimization approach.

The optimization method takes as input the multi-class SVM scores for every frame t of the video. It outputs a sequence $\xi = \{s_t\}$ of camera labels $s_t \in [1 : M]$, where M is the total possible camera labels (five in our case) for all frames $t = \{1 : N\}$. We minimize the following global cost function:

$$E(\xi) = \sum_{t=1}^{N} E_d(s_t) + \sum_{t=2}^{N} E_s(s_{t-1}, s_t). \quad (1)$$

The cost function consists of a data term E_d that measures the evidence of the object state given the SVM scores and a smoothness term E_s which penalizes camera transitions. The data term and the smoothness term are defined as follows:

$$E_d(s_t) = -\log(P(s_t, t)) \ . \quad (2)$$

Here, $P(s_t, t)$ is the SVM classification score for camera label s_t at frame t. And,

$$E_s(s_{t-1}, s_t) = \begin{cases} 0 & \text{if } s_{t-1} = s_t \ , \\ \lambda & \text{otherwise} \ . \end{cases} \quad (3)$$

Where λ is a constant, which is determined empirically. Finally, we use dynamic programming (DP) to solve the optimization problem presented in Equation 1. The DP algorithm constructs a graph with $M \times N$ nodes (M rows, N columns) and computes the minimum cost to reach each node. Finally, we backtrack from the minimum cost node in the last column.

3.2 Event Tagging

Given a video clip, the goal of event tagging is to mark all the salient events on the timeline. It is an important problem in sports analysis as it can produce activity description and other high level results (summarization, highlights generation etc). In our work, we consider four different salient events (Goal, Corner, Substitution, Foul) and a gameplay event (which broadly covers rest of the possible events). The selection was motivated by the online textual commentaries where these four salient events are distinctly marked. An instance for of each of the five classes is illustrated in Figure 4.

Similar to camera label estimation, we design event tagging problem as a per frame classification task. For classifying a frame, the feature vector is built using a total of 40 frames centred at that frame, where each individual frame is represented using the BoW histograms, camera label and motion features. The BoW histograms capture the correlation between the events and the distribution of features, they are obtained in a similar manner as in previous section. The camera label per frame is defined as a five dimensional boolean vector and captures the correlation of an event with camera transitions.

Fig. 4 We classify events into five different categories namely (from left to right) goal, corner, foul, substitution and gameplay.

The motion features represent the correlation of player movements with different events. They are computed from the player tracks obtained using the combination of discriminative trained deformable part models (DPM) proposed by Fezenswalb et al. [5] and the data association approach of Pirsiavash and Ramanan [16]. The player detections are performed per frame basis using a DPM model specifically trained for the case of football videos.

Training DPM: While training DPM's, the choice of negative examples strongly impacted the quality of the detector. Using a fixed set of 200 images with manually annotated bounding boxes for players (2120 instances of players) as positive examples and varying the negative examples, we trained three different versions of DPM:

1. Using randomly sampled windows from frames of football videos as negative examples (not overlapping with players and similar in size). Approx 20000 negative examples were used.
2. Using entire VOC 2007 dataset for negative samples (all classes except pedestrians)
3. Using both VOC 2007 dataset and randomly sampled windows from football videos

We then compared the detection performance of these three cases with the pre trained pedestrian detector on a set of 75 test images. The DPM model trained with only VOC dataset as negative examples gave best results and was finally used for obtaining the player detections.

The per frame detections are then combined into player tracks using [16]. The result is a set of bounding boxes represented by center (x, y) and height(h) and their corresponding labels (representing different tracks). Player tracks shorter than four frames are ignored. Using the corresponding bounding boxes in consecutive frames, we compute a nine dimensional motion feature vector (mean xy-motion; median xy-motion; average h, median h, min-h, max-h and number of corresponding bounding boxes).

The BoW histograms, camera label and motion features are then concatenated for each frame individually. Finally we perform experiments on event classification using both the mean and concatenation of features from 40 individual frames. The classification is performed using a five class SVM which is trained using manually annotated ground truth data. The per frame classification results are then temporally smoothed in a similar way as in Equation 1, penalizing frequent event transitions.

4 Experiments

We perform experiments on broadcast video sequences from 10 matches of football world cup 2014. We present results on the camera label estimation and the salient event classification. We make quantitative comparisons in each case using manually annotated ground truth data. Furthermore, using an application of spatial segmentation, we show that even a simple algorithm with the knowledge of camera viewpoint can bring as much as 20% improvement over the state of the art. Each of these experiments are described with detail in the proceeding sections:

4.1 Camera label estimation

We manually annotated two 45 minutes videos (324000 frames), from two different matches with camera labels (with each frame assigned an unique label) for the quantitative analysis of camera label estimation. One part was used for training and another was used for testing.

We compare our method with two commonly used approaches from the previous work based on color [15, 25] and the motion vectors [24]. The color based approaches classify shots into different classes based on the ratio of the green color or the color histogram analysis. We implemented a similar color based algorithm using ten bin histograms and classified shots into different categories using ratio of the dominant color (we used dominant color instead of a fixed shade of green to bring further robustness). The ratios for each class were learnt from the training data.

For motion vector based classification, we learnt a SVM classifier based on optical flow between consecutive frames. The optical flow was computed by downsampling the image to $p \times p$ pixels, where p denotes the bin size. We tested different bin sizes and only the best result is presented in the paper (using $p = 20$).

	Top-zout	Top-zin	Ground-zout	Ground-zin	Misc
Top-zout	**81.4**	6.9	8.2	3.4	0.1
Top-zin	68.7	**15.9**	9.3	6.1	0
Ground-zout	3	7	**79.7**	10.3	0
Ground-zin	3.3	5.7	70.4	**20.6**	0
Misc	2.5	1.8	33.6	55.7	**6.4**

Table 1 Camera label estimation results using dominant color ratio with five different camera viewpoints (percentages)

	Top-zout	Top-zin	Ground-zout	Ground-zin	Misc
Top-zout	**88.8**	2.3	5.9	2.3	0.7
Top-zin	36.2	**25.9**	17.3	17.7	2.9
Ground-zout	33.6	8.9	**38.4**	16.2	2.9
Ground-zin	28.7	11.6	24.4	**32.4**	2.9
Misc	30	14	22	16	**18**

Table 2 Camera label estimation results using optical flow.

	Top-zout	Top-zin	Ground-zout	Ground-zin	Misc
Top-zout	**98**	0.7	1.2	0.1	0
Top-zin	0.8	**65.9**	8.1	25.1	0.1
Ground-zout	0.4	3.4	**79.6**	15.1	1.5
Ground-zin	0.1	1.4	19	**78.5**	1
Misc	0	0.1	0.2	4	**95.7**

Table 3 Camera label estimation results using our method

We obtained an average accuracy of 45%, 61% and **85**% using color based approach, optical flow based approach and our method respectively using independent per frame classification. The accuracy improved to 56%, 65.4% and **92.2**% respectively, after the MRF smoothing. The confusion matrices of results obtained after MRF smoothing are shown in Table 1, Table 2 and Table 3.

The dominant color ratio based approach (Table 1) fails to do the classification accurately. Strong misclassification between Top-zoom out and Top zoom-in shots can be observed. This occurs due to the large ratio of green color in both top zoom in (usually close up shots from top angle) and top zoom out shots. The color based approach completely fails to classify Miscellaneous (crowd) shots and also heavily confuses among ground zoom-out and ground zoom-in shots. The optical flow based approach also gives noisy results and strongly confuses top zoom out shots with other viewpoints.

On the other hand, our method provides highly accurate results(Table 3). It almost perfectly classifies top zoom out shots which typically cover major part of the football game. The confusion of other classes with top zoom out shots is also negligible (column one in Table 3). In general we give high accuracy for most classes (near 80%) and the confusion when occurs is understandable. For example, in tight close up shots, it is difficult to distinguish between top zoom-in and ground zoom-in camera angles (the background is blurred in both cases and only the player profile contributes to the features).

	Corner	Foul	Goal	Gameplay	Subst.
Corner	**43.4**	4.9	31.3	9.8	10.6
Foul	0	**48.5**	32.9	13.7	4.9
Goal	2.5	8.9	**82**	4.5	2.1
Gameplay	3.9	4.2	5.8	**85.2**	0.9
Subst.	1.6	25.5	41.1	0.7	**31.1**

Table 4 Event tagging results using only the BoW histograms considering five different events (percentages).

	Corner	Foul	Goal	Gameplay	Subst.
Corner	**71.2**	0	28.8	0	0
Foul	0.5	**43.8**	40.2	8	7.5
Goal	0.6	8.7	**81.2**	4.2	5.3
Gameplay	3.4	3	2.4	**91.2**	0
Subst.	0	12.7	32.7	0	**54.6**

Table 5 Event tagging results using the combination of camera label information with the BoW histograms.

	Corner	Foul	Goal	Gameplay	Subst.
Corner	**75.1**	5.4	18	0.6	0.9
Foul	0.2	**57.4**	22	10.4	10
Goal	2.7	3.4	**84.6**	4.2	5.1
Gameplay	0.9	1.5	1.5	**94.6**	1.5
Subst.	1.3	9.6	24	2.7	**62.4**

Table 6 Event tagging results using the combination of BoW histograms; camera label information and the motion features.

4.2 Event tagging

For event tagging experiment, we created a dataset of 176 clips encompassing all the five events. The dataset includes 43 goal sequences, 30 corner sequences, 42 substitution sequences, 27 fouls sequences and 34 gameplay sequences. The clips were extracted from the video by manually annotating the start and the end of the event. The length of each sequence is different and is defined based on the knowledge of the game. For example, a goal event starts from the point the ball enters the goal post and ends at the new kickoff (so the goal event includes all the celebrations). Similarly, a corner event starts just before the corner kick and ends at the first deflection. The sequences for gameplay event were of random length. We used 60% of the data for training/validation and 40% for testing. Example videos of each kind of event with qualitative results are provided in the supplementary material. Another qualitative result on an interesting 30 minutes sequence (with five goals in quick succession) is shown in Figure 1.

For quantitative analysis, we sampled test data clips from different classes and joined them in random order to create a large video sequence. The event classification task was then performed per frame basis on this large sequence. We tested classification task using both the mean (taking mean of features from 40 frames) and concatenation (concatenating features from 40 frames). The average accuracy of around 80% was obtained in both the cases, which improved to 85.5% after MRF smoothing.

The confusion matrix of results after MRF smoothing (using mean features from 40 frames) is given in Table 6. The goal event and gameplay event are predicted with an accuracy over 85%. The confusion when occurs, is understandable, for example some corner events lead into goals or near misses and in such instances the corner event is misclassified as goal event. Similarly substitution event is often followed with the crowd viewpoint (cheering the player) and replay of goals (or assists) by substituted player, which leads to confusion between substitution and goal event.

Comparing results in Table 4, Table 5 and Table 6 we can observe that both context(camera label information) and motion features bring significant improvements in event recognition task. For example, the accuracy of predicting corner improves by almost 30% after including camera label information (Table 4 and Table 5). Similarly, including motion information brings almost 15% improvement in predicting Foul event (Table 5 and Table 6). Overall, the average accuracies increased from 77% to 81.8% after adding camera context. The average accuracy further improved to 85.5% after the inclusion of motion features. In terms of time complexity, the initial method using only BoW histograms runs at around 1 fps in our current implementation on a single core intel-i5 CPU with 16GB memory. Adding camera label information does not bring any significant change, but including motions cues increases computational time to around 0.4 fps.

We further compare our method with a hidden conditional random field (HCRF) based method [17] and a recent action recognition using trajectory-pooled-deep-convolutional descriptors [21] (with combined spatial and temporal features). For comparison with [17] we use our own implementation for computing features with sequence classification toolbox from [20]. The results with precision for predicting each event is presented in Table 7. We can clearly observe that our approach outperforms the generic state of the art action recognition method [21], which shows the importance of the context. Our method also improves the precision over a more sophisticated(using ad-hoc domain specific features) sequence classification method [21] on most of the events.

	Corner	Foul	Goal	Gameplay	Subst.
TDD	21.05	12.82	77.45	77.08	25.80
HCRF	7.6	15.5	6.6	77.13	**86.44**
Ours	**75.1**	**57.4**	**84.6**	**94.6**	62.4

Table 7 Comparison of our method with TDD [21] and HCRF [17].

4.3 Spatial segmentation

In this section, we demonstrate the usefulness of camera label information with an application of spatial segmentation. Given an individual frame, the task is to assign each pixel with an unique class label. We consider three different classes i.e players; crowd and playing field. The experiments are performed on a dataset of 50 images with top-zoom out camera labels and 50 images with zoom-in camera labels (sampled both from ground zoom-in and top zoom-in cameras). The images are equally sampled from 10 different matches to cover different challenges in segmentation like shadows, different color player jerseys etc. The ground truth labels for all the 100 images are created manually.

We investigate two different segmentation approaches. First we follow the class based image segmentation approach of Ladicky et al. [12] which casts the segmentation problem as graph cut based inference on conditional random fields. Second we propose a segmentation approach specific to top zoom-out views. Our main insight is that knowing the context can help to design minimal segmentation algorithms bringing significant improvement in terms of both time and labelling accuracy. Knowing that the camera label is a top zoom-out, we can assume that the large part of the frame will be covered by playing field which can be efficiently segmented based on color. We use a variation of Heckbert's [6] median cut algorithm to estimate dominant color and segment out the playing field in normalized rgb space. The segmentation is performed using a threshold on euclidean distance from the dominant color. The holes in the playing field are then labelled as players and the rest is labelled as crowd.

We then trained three instances of automatic labelling environment (ALE) [12], first for the top zoom-out camera labels, second for the zoom-in camera labels and third for the combined set. Half of the total images were used for training in each case. The ALE segmentation results for top zoom-out and zoom-in are illustrated in Table 8. We can observe that ALE fails to segments players accurately in images with top zoom-out camera labels. Interestingly when we trained ALE by taking frames with from both zoom-out and zoom-in viewpoints we obtained nearly same results. This suggests that ALE is not taking full advantage of the context. We then performed segmentation using the second context aware approach (based on dominant color) and the recall for the class *Players* improved to **64.2**% maintaining nearly the same recall for *Field* and *Crowd* classes. The results from the dominant color based approach were obtained in real time compared to 30 hours training (for 50 images with image resolution of 720p)

	ALE top zoom-out	ALE zoom-in
Players	44.4	80.8
Field	99.31	93.3
Crowd	99.8	91.5

Table 8 Results of average per class recall measure, defined as $\frac{TruePositives}{TruePositives+FalsePositives}$ on ALE [12]. The recall measure for class *Players* is low in top zoom-out viewpoints.

and 3 hours testing (4-5 minutes per frame) in case of ALE. This clearly shows that knowing context (like camera labels) can bring significant improvements for the task of spatial segmentation both in terms of performance and recall.

5 Conclusions

In this paper we have investigated the problem of automatic analysis of football broadcast videos. We have shown that this problem can be partitioned into two smaller problems, namely, camera viewpoint estimation and event tagging. We have demonstrated that since the input videos are already edited, the camera viewpoint information provides a natural context which could be exploited to improve the other task of event tagging.

Based on thorough quantitative analysis on variety of tasks in 10 football matches, we have justified our claims. Our method obtains an overall accuracy of 92.2% for camera viewpoint estimation and 85.5% accuracy for event classification. We have also demonstrated that the contextual approach can outperform state of the art deep learning based action recognition approaches. We further demonstrate that the accuracy of tasks like spatial semantic segmentation can be improved by as much as 20% using the context.

References

1. Assfalg, J., Bertini, M., Colombo, C., Bimbo, A.D., Nunziati, W.: Semantic annotation of soccer videos: automatic highlights identification. CVIU **92**, 285 – 305 (2003)
2. Chen, C., Wang, O., Heinzle, S., Carr, P., Smolic, A., Gross, M.: Computational sports broadcasting: Automated director assistance for live sports. In: ICME (2013)
3. Csurka, G., Dance, C., Fan, L., Willamowski, J., Bray, C.: Visual categorization with bags of keypoints. In: Workshop on statistical learning in computer vision, ECCV, vol. 1, pp. 1–2 (2004)
4. Duan, L.Y., Xu, M., Tian, Q., Xu, C.S., Jin, J.S.: A unified framework for semantic shot classification in sports video. Multimedia, IEEE Transactions on **7**(6), 1066–1083 (2005)
5. Felzenszwalb, P.F., Girshick, R.B., McAllester, D., Ramanan, D.: Object detection with discriminatively trained part based models. TPAMI **32**(9), 1627–1645 (2010)
6. Heckbert, P.: Color image quantization for frame buffer display. In: SIGGRAPH (1982)
7. Jain, V., Singhal, A., Luo, J.: Selective hidden random fields: Exploiting domain-specific saliency for event classification. In: CVPR (2008)
8. Kapela, R., McGuinness, K., Swietlicka, A., OConnor, N.E.: Real-time event detection in field sport videos. In: Computer Vision in Sports, pp. 293–316 (2014)
9. Kim, K., Grundmann, M., Shamir, A., Matthews, I., Hodgins, J., Essa, I.: Motion fields to predict play evolution in dynamic sport scenes. In: CVPR (2010)
10. Kong, Y., Hu, W., Zhang, X., Wang, H., Jia, Y.: Learning group activity in soccer videos from local motion. In: ACCV (2010)
11. Koprinska, I., Carrato, S.: Temporal video segmentation: A survey. Signal processing: Image communication **16**(5), 477–500 (2001)
12. Ladicky, L., Russell, C., Kohli, P., Torr, P.H.: Graph cut based inference with co-occurrence statistics. In: ECCV (2010)
13. Lucey, P., Bialkowski, A., Carr, P., Morgan, S., Matthews, I., Sheikh, Y.: Representing and discovering adversarial team behaviors using player roles. In: CVPR (2013)
14. Ma, Z., Yang, Y., Cai, Y., Sebe, N., Hauptmann, A.G.: Knowledge adaptation for ad hoc multimedia event detection with few exemplars. In: ACM Multimedia, pp. 469–478 (2012)
15. Nguyen, N., Yoshitaka, A.: Shot type and replay detection for soccer video parsing. In: Multimedia (ISM), 2012 IEEE International Symposium on, pp. 344–347 (2012)
16. Pirsiavash, H., Ramanan, D., Fowlkes, C.C.: Globally-optimal greedy algorithms for tracking a variable number of objects. In: CVPR (2011)
17. Qian, X., Liu, G., Wang, Z., Li, Z., Wang, H.: Highlight events detection in soccer video using hcrf. In: Proceedings of the Second International Conference on Internet Multimedia Computing and Service, pp. 171–174 (2010)
18. Sigari, M.H., Soltanian-Zadeh, H., Kiani, V., Pourreza, A.R.: Counterattack detection in broadcast soccer videos using camera motion estimation. In: AISP, pp. 101–106 (2015)
19. Thompson, R., Bowen, C.: Grammar of the Edit. Focal Press (2009)
20. Walecki, R., Rudovic, O., Pavlovic, V., Pantic, M.: Variable-state latent conditional random fields for facial expression recognition and action unit detection. In: Automatic Face and Gesture Recognition (2015)
21. Wang, L., Qiao, Y., Tang, X.: Action recognition with trajectory-pooled deep-convolutional descriptors. In: CVPR (2015)
22. Xie, L., Chang, S.F., Divakaran, A., Sun, H.: Structure analysis of soccer video with hidden markov models. In: ICASSP (2002)
23. Xu, C., Wang, J., Wan, K., Li, Y., Duan, L.: Live sports event detection based on broadcast video and web-casting text. In: ACM Multimedia, pp. 221–230 (2006)
24. Xu, G., Ma, Y.F., Zhang, H.J., Yang, S.Q.: An hmm-based framework for video semantic analysis. Circuits and Systems for Video Technology, IEEE Transactions on **15**(11), 1422–1433 (2005)
25. Xu, P., Xie, L., Chang, S.F., Divakaran, A., Vetro, A., Sun, H.: Algorithms and system for segmentation and structure analysis in soccer video. In: ICME, vol. 1, pp. 928–931 (2001)

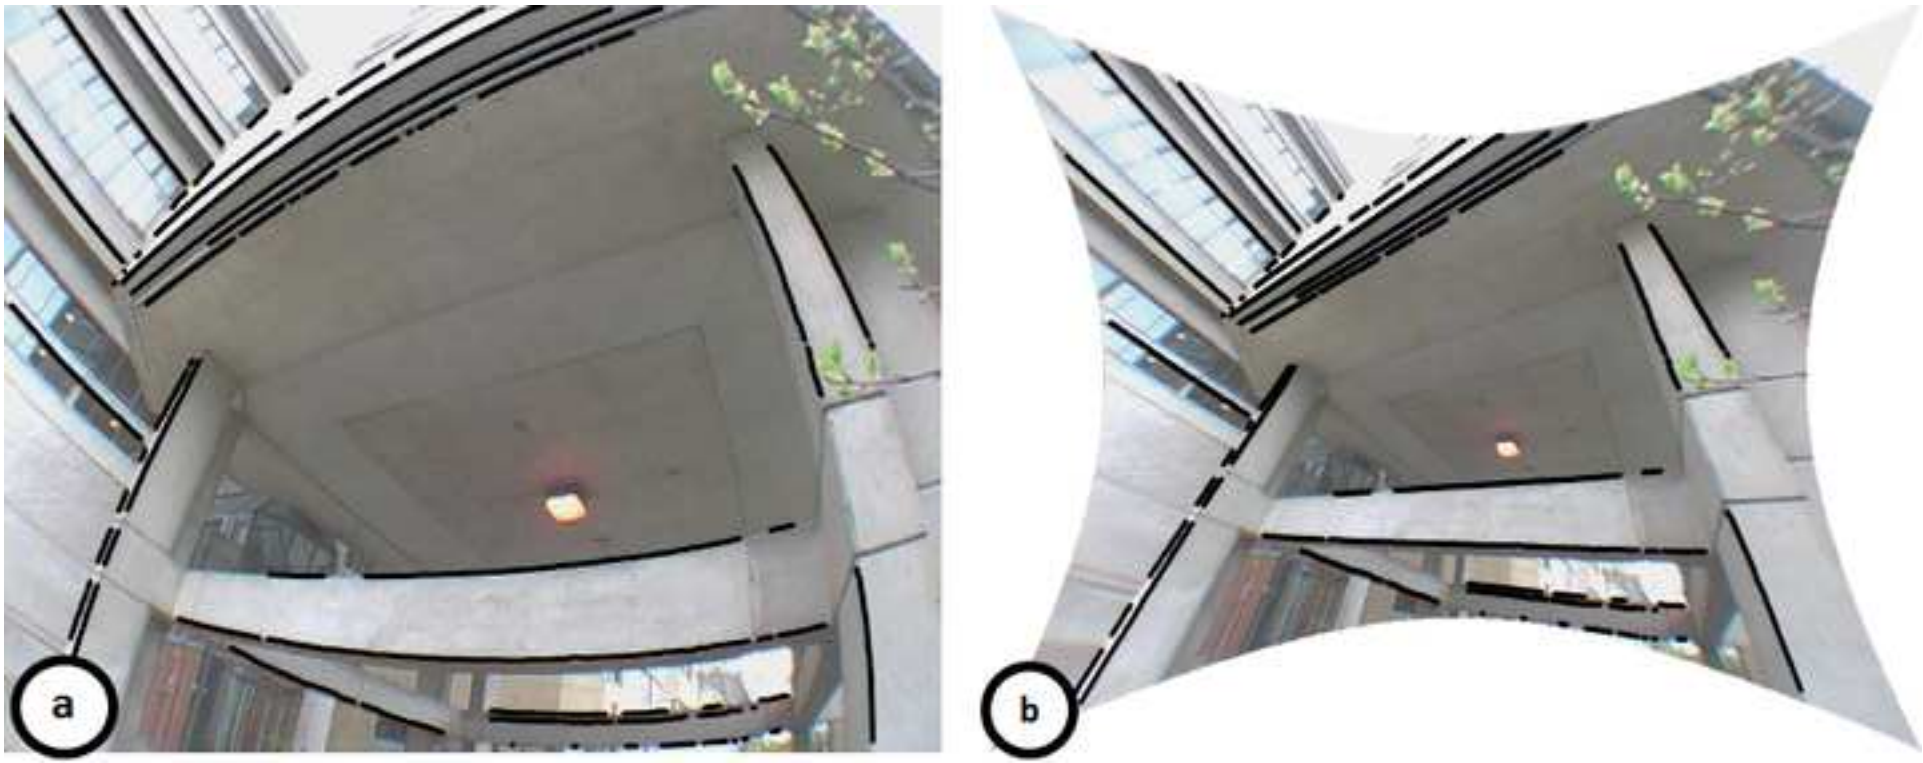

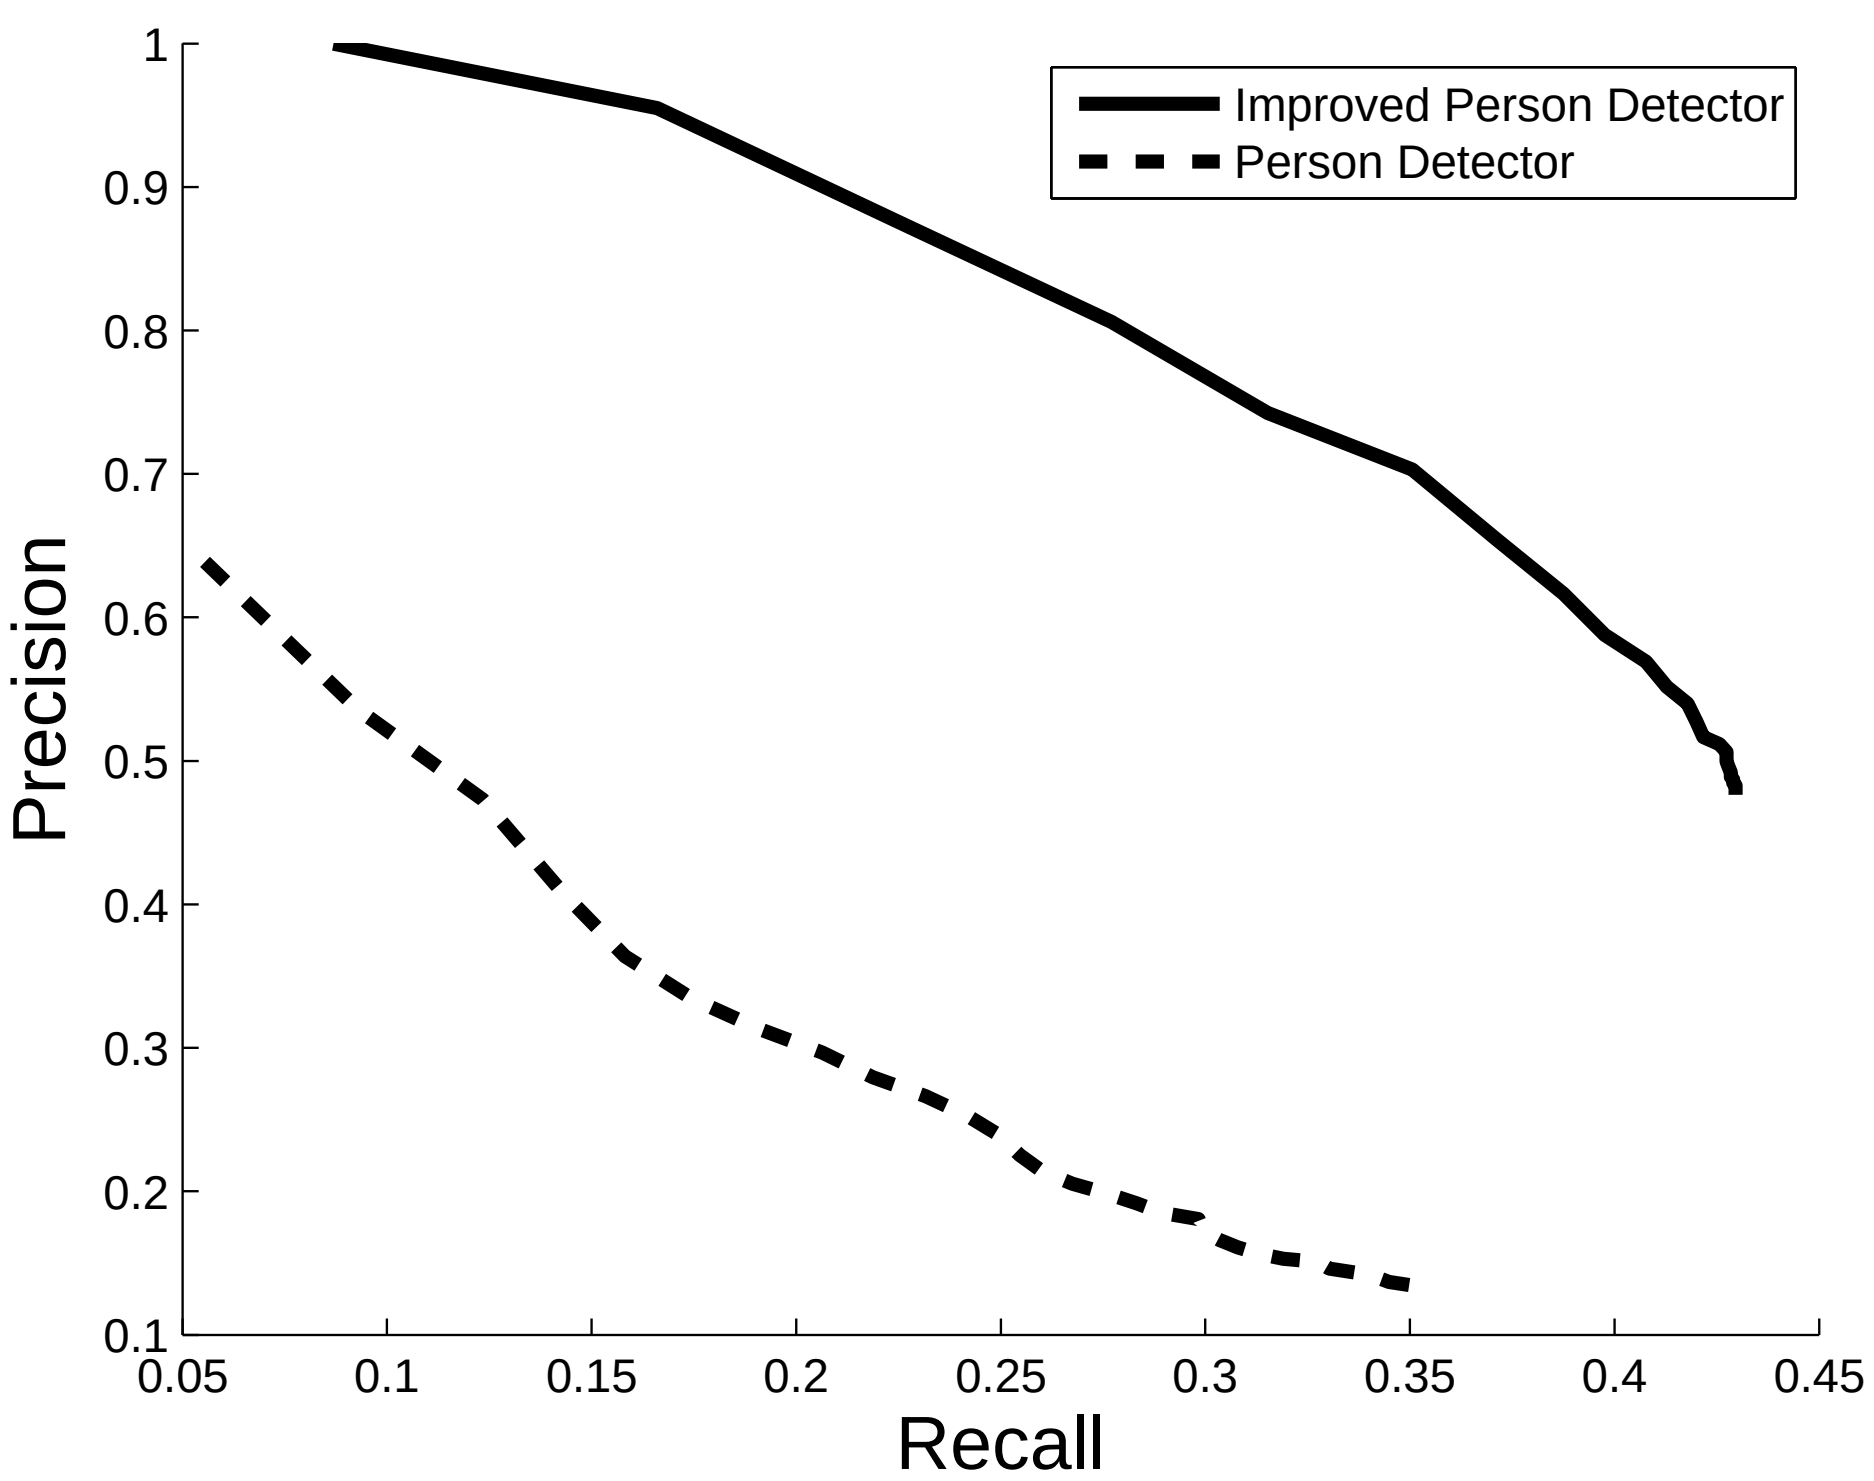

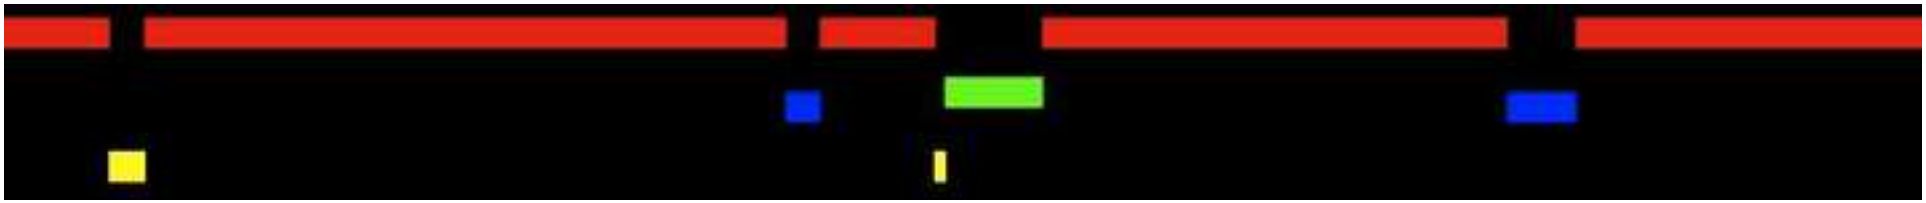

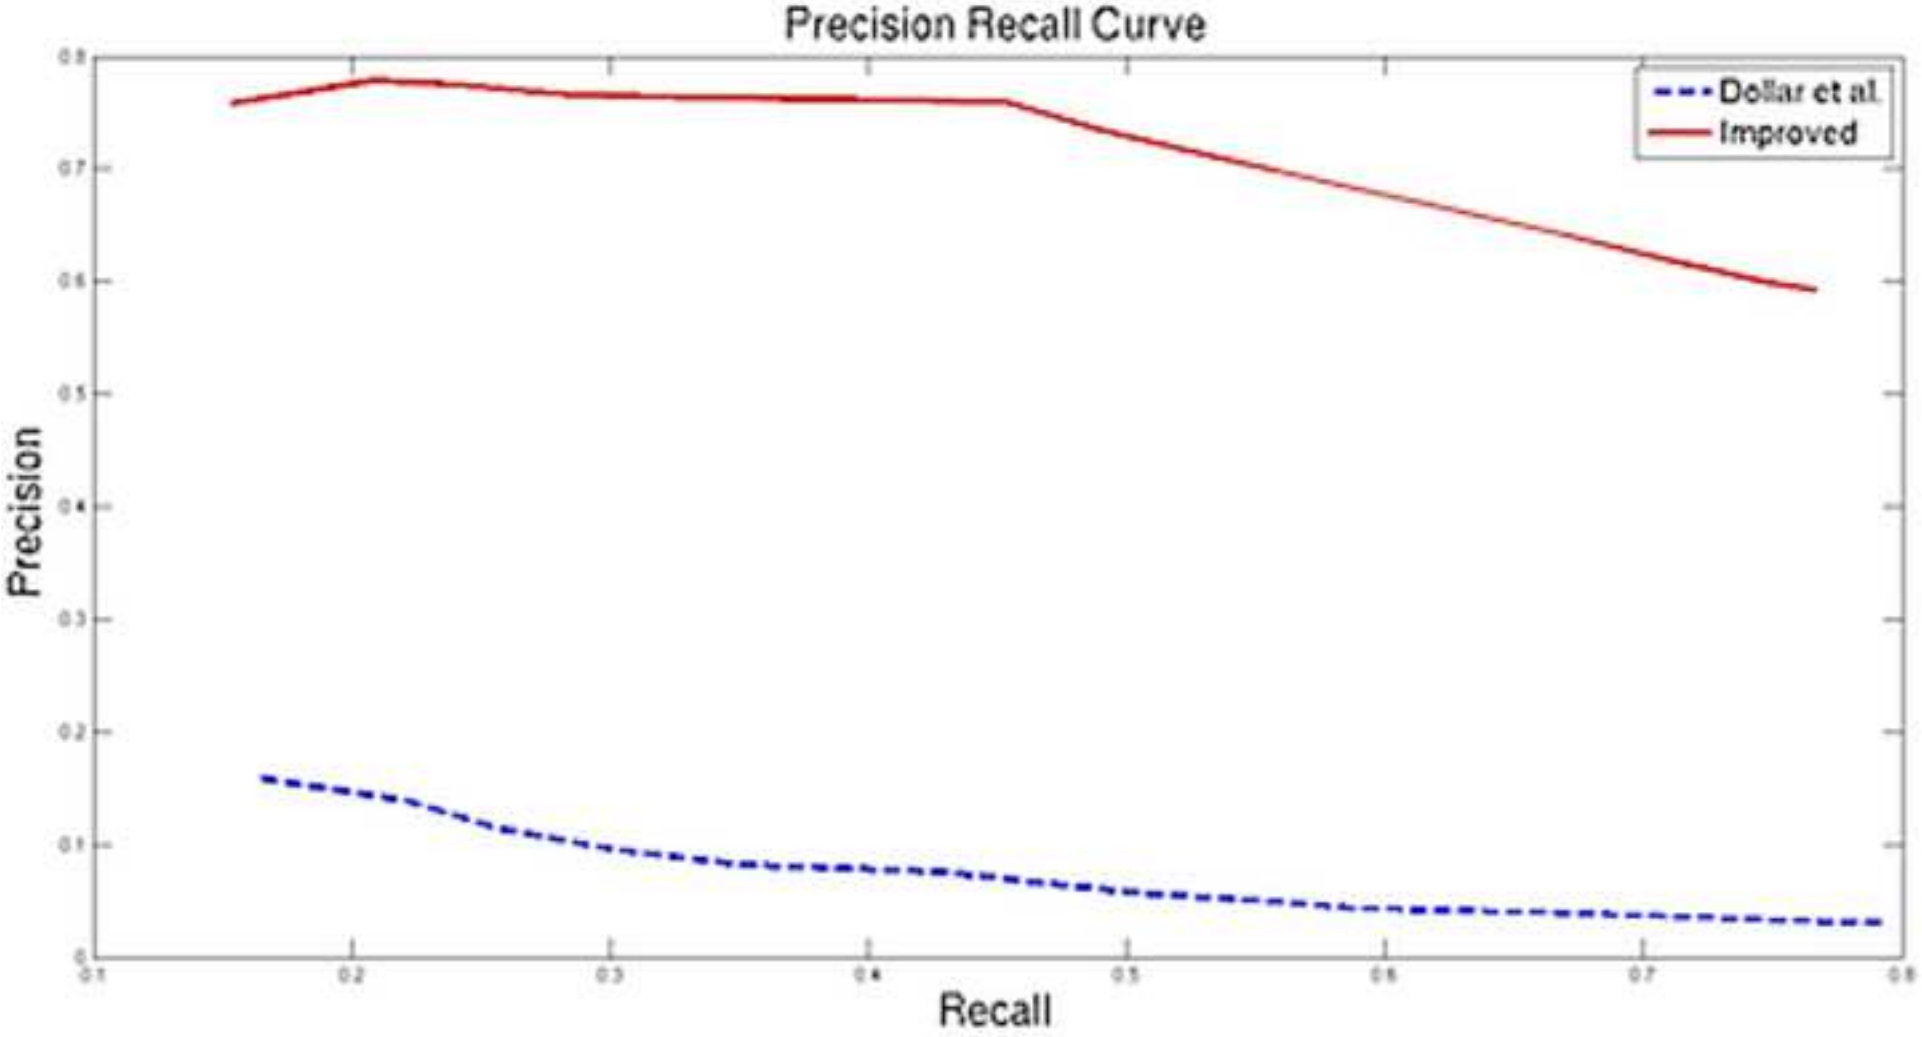

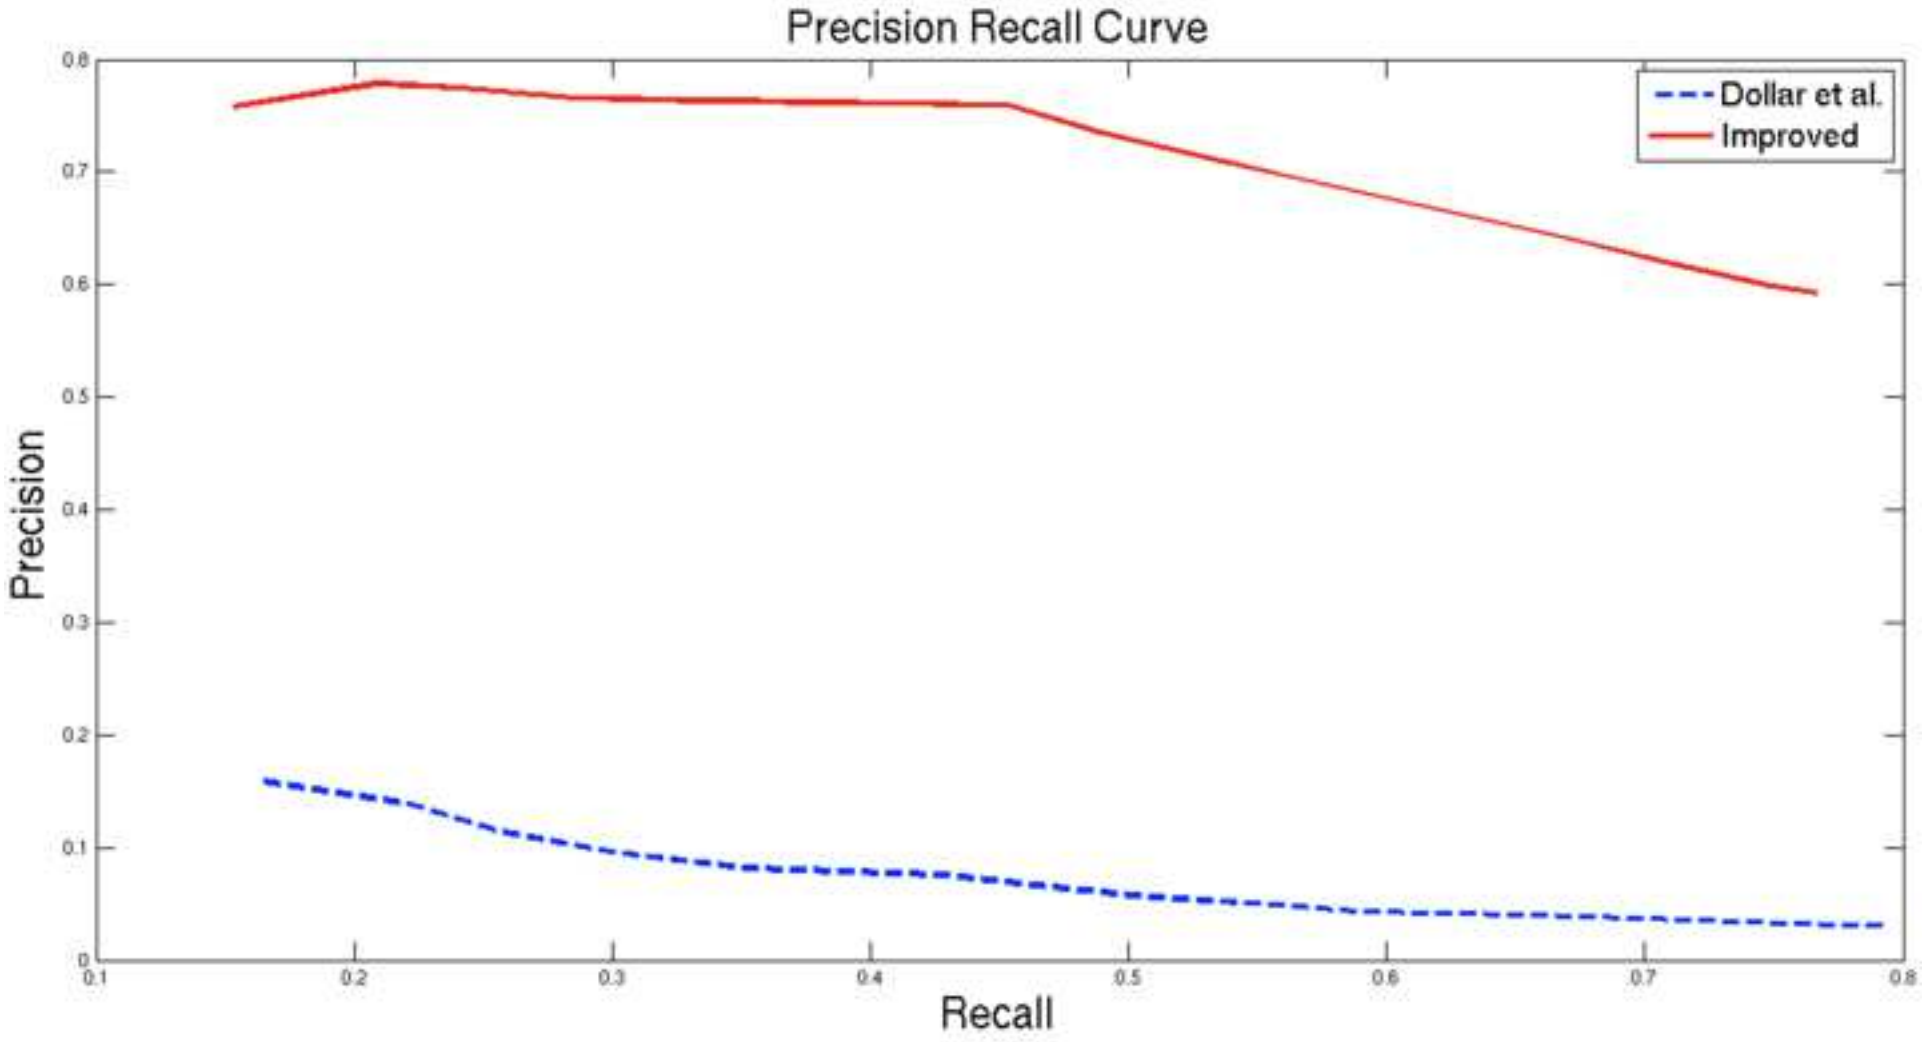

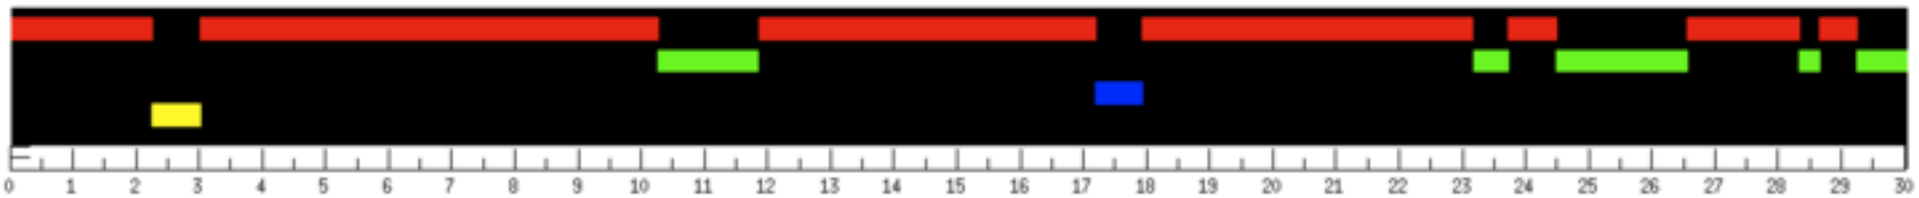

Click here to view linked References

Click here to view linked References

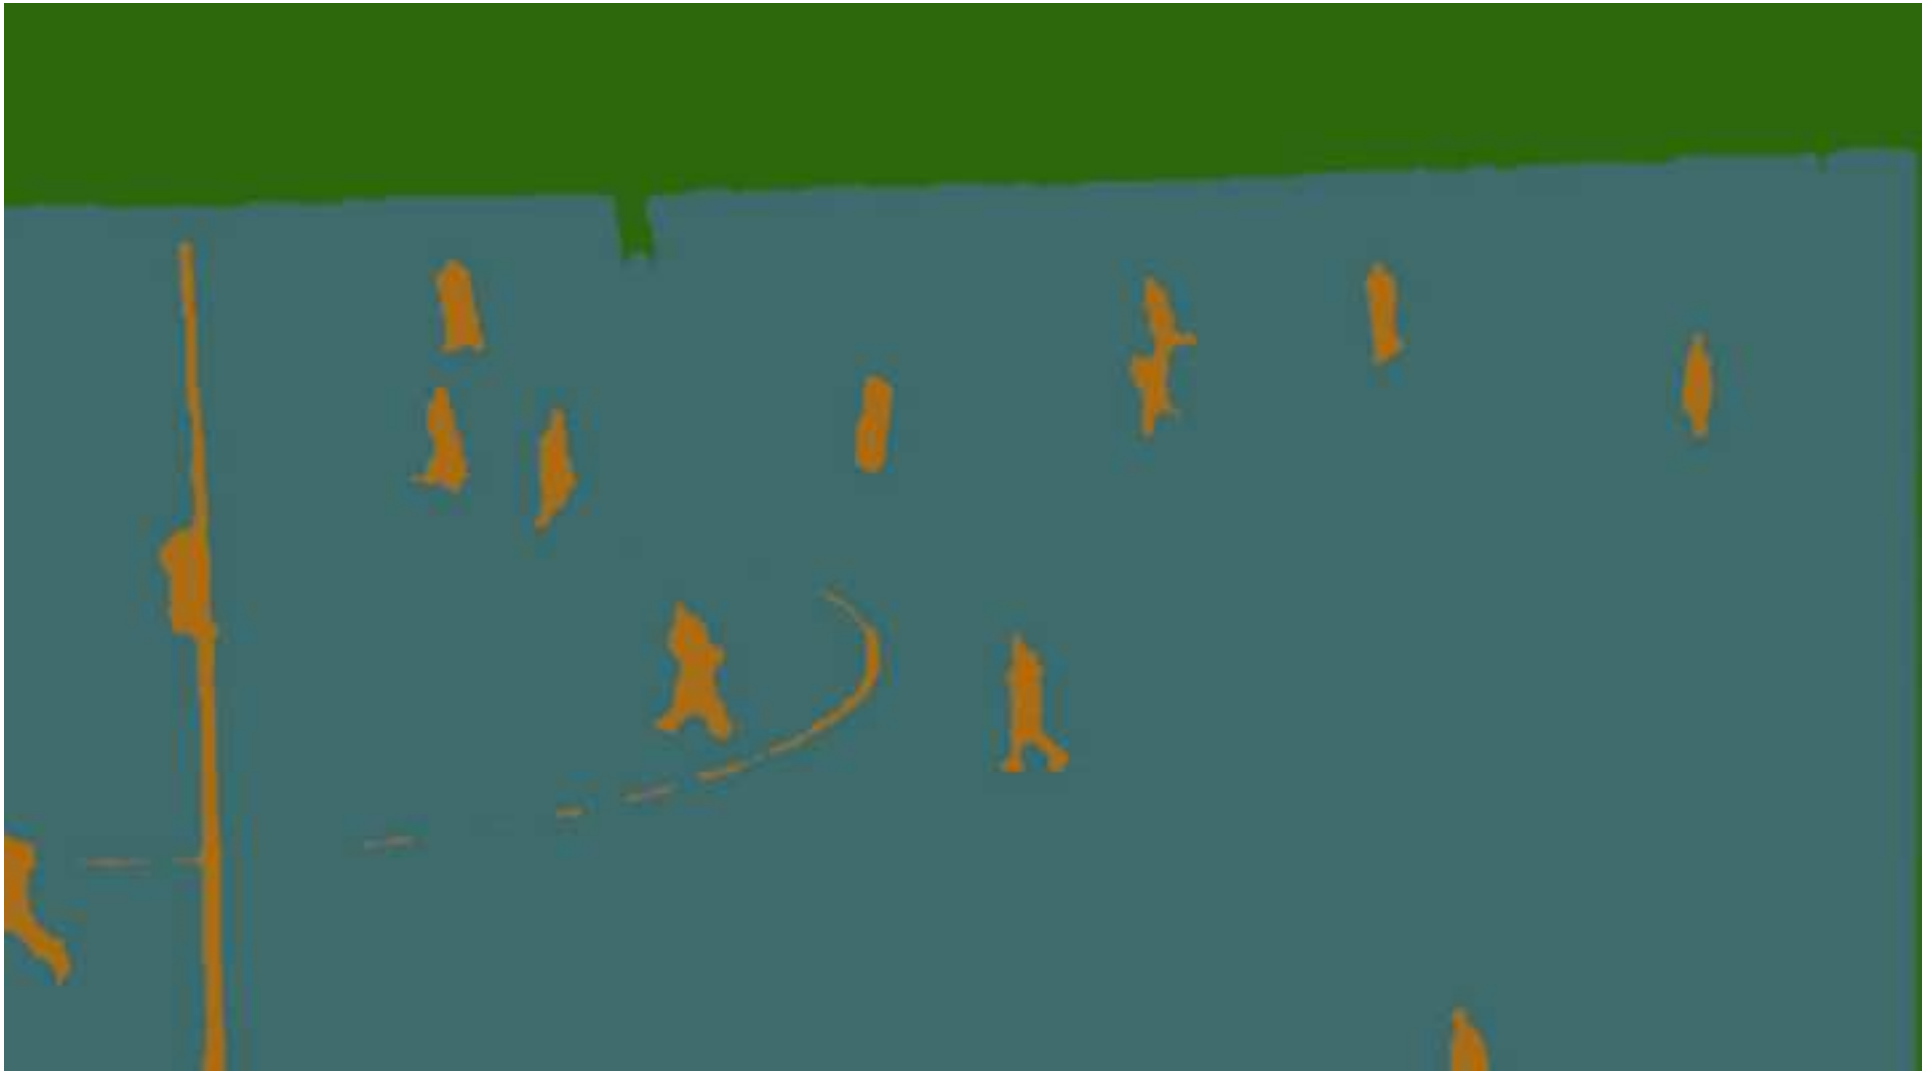

Click here to view linked References

Click here to view linked References

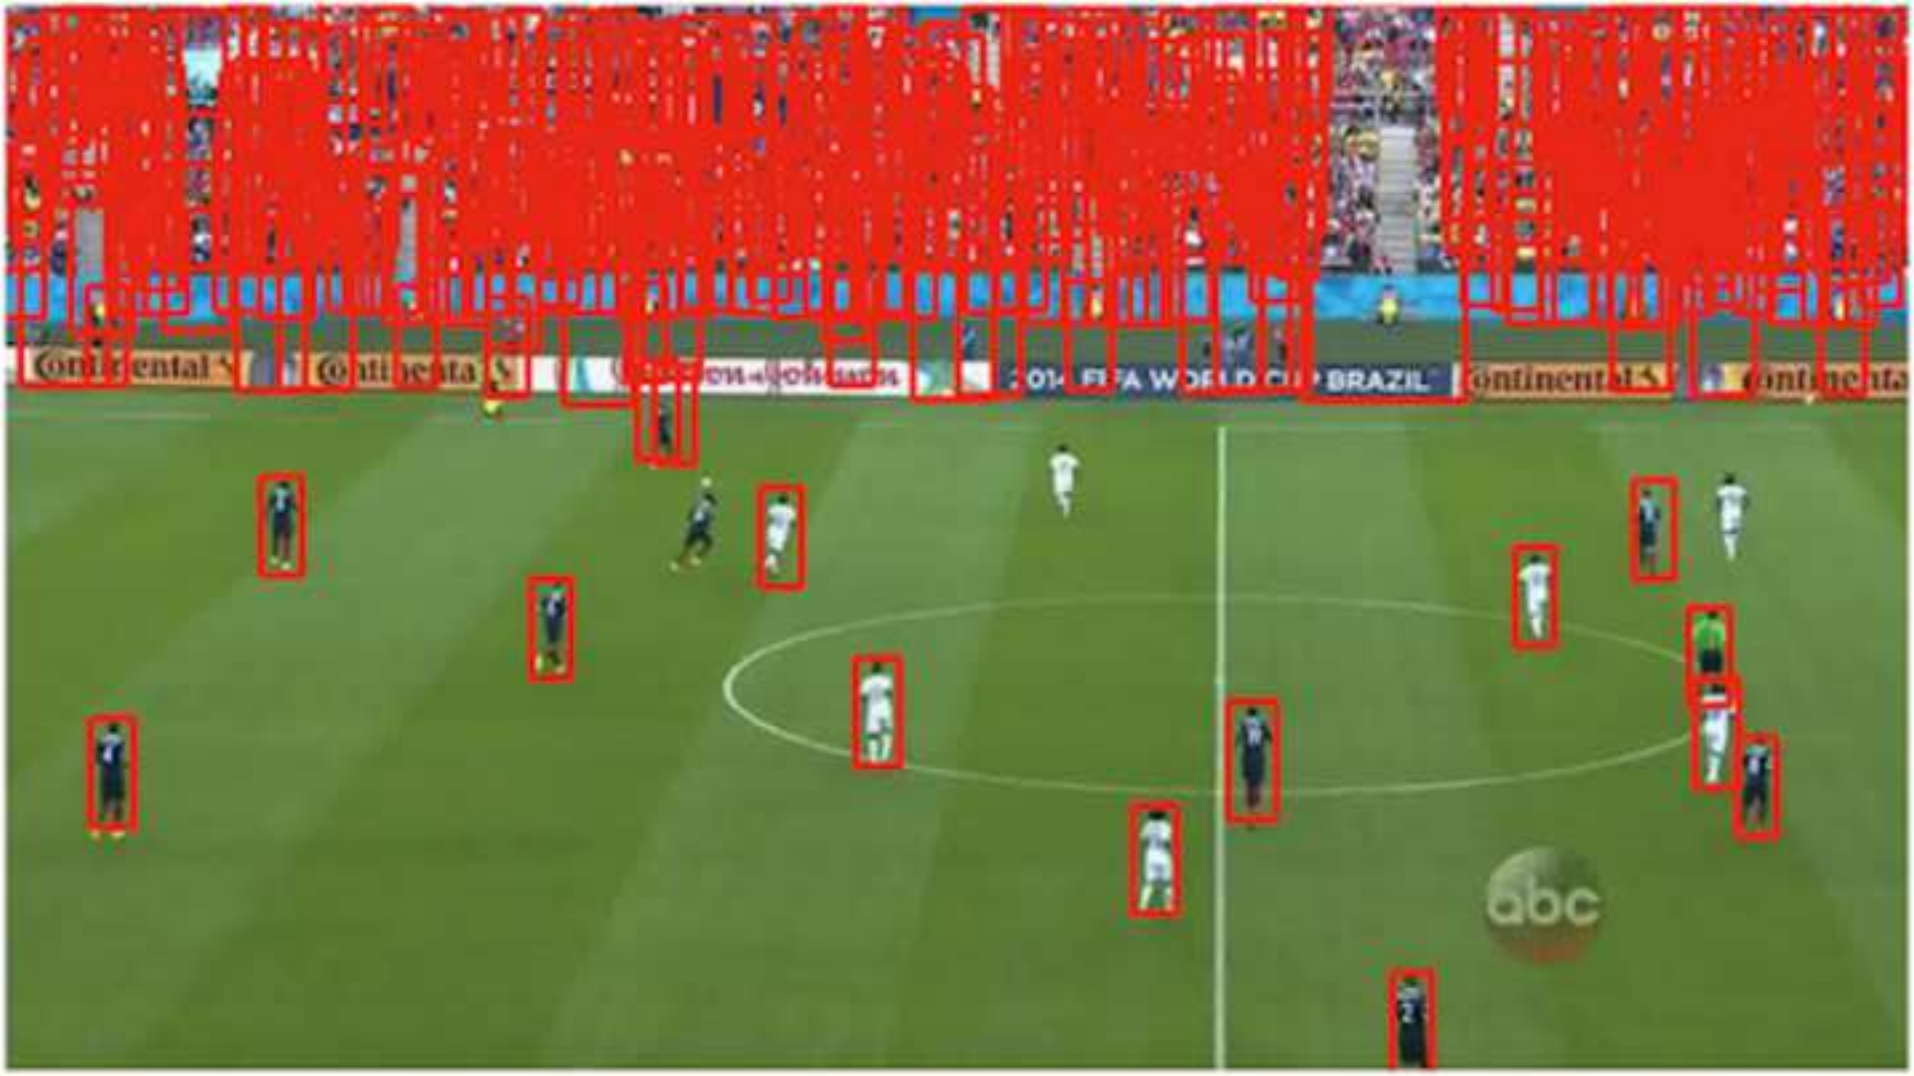

Click here to view linked References

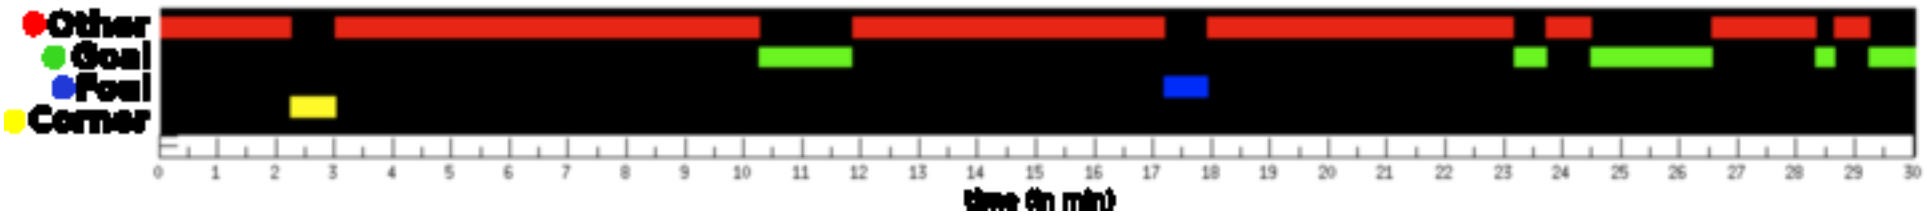

Click here to view linked References

Ground Zoomin	Ground ZoomOut	Top Zoomin	Top ZoomOut	Top ZoomOut_2	Crowd Zoomin	Crowd ZoomOut
						

Example qualitative result

Example qualitative result

Example qualitative result

Example qualitative result

Example qualitative result